

CGS698C - Assignment 4

Sahil Tomar (210898)

2024-07-03

Part 1 : Power posing and testosterone

```
df_powerpose <- read.table("df_powerpose.csv", header=T, sep=",")
flextable(head(df_powerpose))
```

	X	id	hptreat	female	age	testm1	testm2
	2	29	High	Male	19	38.725	62.375
	3	30	Low	Female	20	32.770	29.235
	4	31	High	Female	20	32.320	27.510
	5	32	Low	Female	18	17.995	28.655
	7	34	Low	Female	21	73.580	44.670
	8	35	High	Female	20	80.695	105.485

The data set shows the testosterone levels of 39 different individuals, before (testm1) and after (testm2) treatment, where treatment refers to each individual being assigned to a high power pose or a low power pose (hptreat) .

The research hypothesis is that on average, assigning a subject a high power pose vs. a low power pose will lead to higher testosterone levels after treatment.

```
df_powerpose <- df_powerpose %>%
  mutate(
    hptreat_binary = ifelse(hptreat == "High", 1, 0),
    change_testosterone = testm2 - testm1
  )

mfit <- brm(
  formula = change_testosterone ~ hptreat_binary,
  data = df_powerpose,
  family = gaussian(),
  chains = 4, cores = 4,
  iter = 4000, warmup = 1000,
  verbose = FALSE
)
```

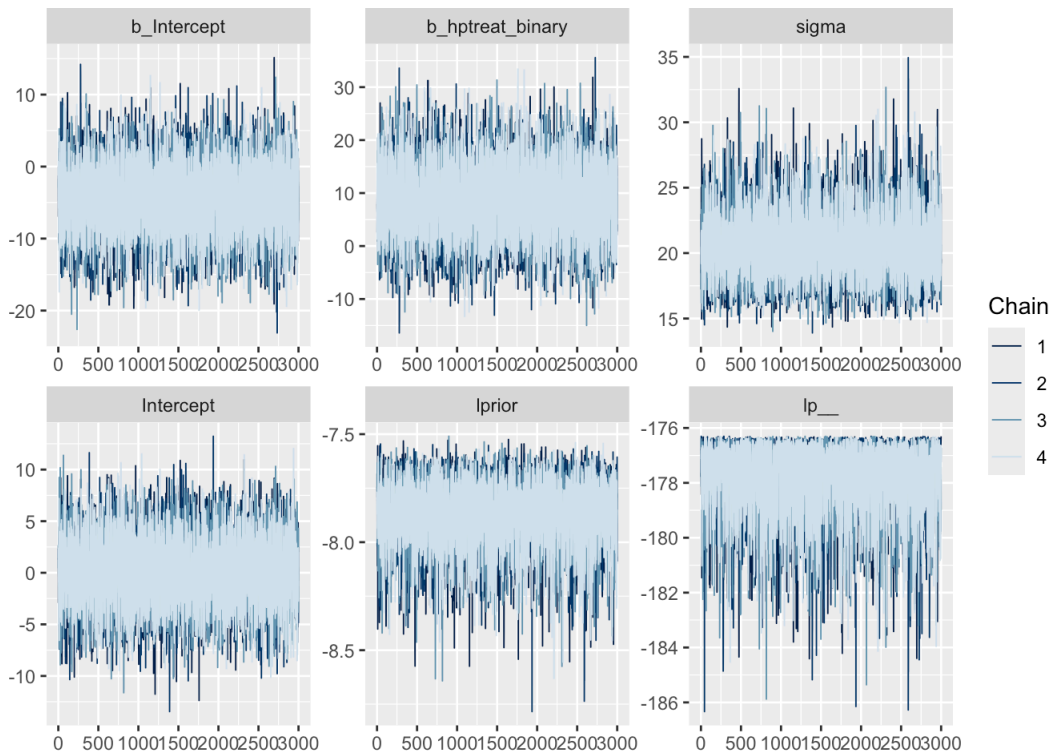
```
summary(mfit)
```

```
## Family: gaussian
## Links: mu = identity; sigma = identity
## Formula: change_testosterone ~ hptreat_binary
## Data: df_powerpose (Number of observations: 39)
## Draws: 4 chains, each with iter = 4000; warmup = 1000; thin = 1;
## total post-warmup draws = 12000
##
## Regression Coefficients:
##           Estimate Est.Error l-95% CI u-95% CI Rhat Bulk_ESS Tail_ESS
## Intercept    -4.49      4.68   -13.79    4.50 1.00   11145    8478
## hptreat_binary  8.89      6.59    -3.99   21.66 1.00   11760    8564
##
## Further Distributional Parameters:
##           Estimate Est.Error l-95% CI u-95% CI Rhat Bulk_ESS Tail_ESS
## sigma    20.55      2.41   16.46   25.83 1.00    9095    8312
##
## Draws were sampled using sampling(NUTS). For each parameter, Bulk_ESS
## and Tail_ESS are effective sample size measures, and Rhat is the potential
## scale reduction factor on split chains (at convergence, Rhat = 1).
```

The above summary provides the value of the `hptreat_binary` coefficient (β) which is the measure of the effect of the variable that encodes the change in testosterone.


The α is mostly in the negative region while β is mostly in the positive regions, this suggests that testosterone tends to increase after a high power pose.

```
mcmc_trace(mfit)
```



```
mcmc_hist(mfit, pars=c("b_Intercept", "b_hptreat_binary", "sigma"))
```

```
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```



Prior Sensitivity Analysis

Using the above summary knowledge, these priors are some priors one can check for sensitivity -

```

# Define prior specifications for sensitivity analysis
prior_specs <- list(
  list(alpha = "normal(-5, 1)", beta = "normal(8, 1)"),
  list(alpha = "normal(-4, 1)", beta = "normal(8, 1)"),
  list(alpha = "normal(-5, 2)", beta = "normal(8, 1)"),
  list(alpha = "normal(-4, 2)", beta = "normal(8, 1)"),
  list(alpha = "normal(-5, 1)", beta = "normal(9, 1)"),
  list(alpha = "normal(-4, 1)", beta = "normal(9, 1)"),
  list(alpha = "normal(-5, 2)", beta = "normal(9, 1)"),
  list(alpha = "normal(-4, 2)", beta = "normal(9, 1)")
)

# Initialize an empty data frame to store sensitivity results
df.sensitivity <- data.frame(matrix(nrow = 0, ncol = 5))
colnames(df.sensitivity) <- c("prior_alpha", "prior_beta", "mean_beta", "q.lower", "q.upper")

# Fit the model and store results for each prior specification
for (spec in prior_specs) {
  # Update priors for alpha and beta
  priors <- c(
    set_prior(spec$alpha, class = "Intercept"),
    set_prior(spec$beta, class = "b", coef = "hptreat_binary"),
    set_prior("normal(0, 10)", class = "sigma") # Keep the prior for sigma unchanged
  )

  # Fit the model
  mfit <- brm(
    formula = change_testosterone ~ hptreat_binary,
    data = df_powerpose,
    family = gaussian(),
    prior = priors,
    chains = 4,
    cores = 4,
    iter = 2000,
    warmup = 1000,
    seed = 123 # Set a seed for reproducibility
  )

  # Extract posterior samples of beta (hptreat_binary)
  post_samples <- posterior_samples(mfit)

  # Calculate mean, lower, and upper quantiles of beta (hptreat_binary)
  mean_beta <- mean(post_samples$b_hptreat_binary)
  lower_quantile <- quantile(post_samples$b_hptreat_binary, probs = 0.025)
  upper_quantile <- quantile(post_samples$b_hptreat_binary, probs = 0.975)

  # Store results in df.sensitivity
  df.sensitivity[nrow(df.sensitivity) + 1, ] <- c(
    spec$alpha,
    spec$beta,
    mean_beta,
    lower_quantile,
    upper_quantile
  )

  # Save the model for each prior specification if needed
}

```

```
flextable(df.sensitivity)
```

prior_alpha	prior_beta	mean_beta	q.lower	q.upper
normal(-5, 1)	normal(8, 1)	8.03562027075994	6.07380653329301	9.99240813270475
normal(-4, 1)	normal(8, 1)	8.02379996721376	6.11090050841724	9.91721994045697
normal(-5, 1)	normal(8, 1)	8.02669329520304	6.11467300835655	9.9432372410656

prior_alpha	prior_beta	mean_beta	q.lower	q.upper
2)	1)			
normal(-4, 2)	normal(8, 1)	8.02424678426533	6.0597824934006	9.90882312485173
normal(-5, 1)	normal(9, 1)	8.98046596349485	7.06759186473789	10.8960662359089
normal(-4, 1)	normal(9, 1)	8.99863379987108	7.09718310024839	10.8531113415611
normal(-5, 2)	normal(9, 1)	8.97507838553259	7.05452706582347	10.9508412809515
normal(-4, 2)	normal(9, 1)	9.00312282250039	7.0918239009274	10.8945481320449

Part 2 : Poisson regression models and hypothesis testing

Exercise 2.1 -

Implement the model in R or Python such that the function gives the number of crossings as the outcome, and takes sentence length, α , and β as its arguments.

We are given:

- The number of crossing dependencies in a sentence can be given by a Poisson distribution - $N_i \sim \text{Poisson}(\lambda_i)$
- where N_i is the number of crossing dependencies in the sentence i
- λ_i is rate parameter indicating the expected rate of crossing dependencies in the sentence i , such that $\log \lambda_i = \alpha + \beta L_i$
- where L_i is the length of the sentence i , α is the expected rate of crossings in a sentence of average length and β is the change in rate of crossings as a function of sentence length.

```
calculate_crossings <- function(sentence_length, alpha, beta) {
  lambda_i <- exp(alpha + beta * sentence_length)
  crossings <- rpois(1, lambda_i)
  return(crossings)
}
```

Above function can be used to get the number of crossings as the outcome given the sentence length, α , and β as its arguments.

Exercise 2.2 -

Generate prior predictions of the model for sentences of length 4 under the following prior assumptions -

$\alpha \sim \text{Normal}_{lb=0}(0.15, 0.1)$ and $\beta \sim \text{Normal}_{lb=0}(0.25, 0.05)$

```
# Sample alpha from truncated Normal(0.15, 0.1)
alpha <- rtruncnorm(1000, a = 0, b = Inf, mean = 0.15, sd = 0.1)

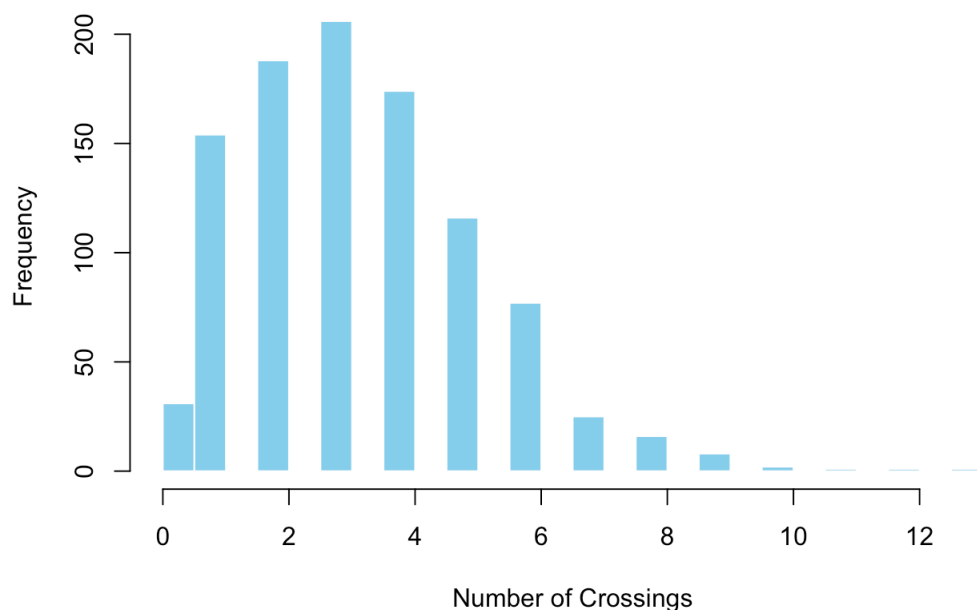
# Sample beta from truncated Normal(0.25, 0.05)
beta <- rtruncnorm(1000, a = 0, b = Inf, mean = 0.25, sd = 0.05)

sentence_length <- 4

prior_crossings <- numeric(length(alpha))
for (i in 1:length(alpha)) {
  prior_crossings[i] <- calculate_crossings(sentence_length, alpha[i], beta[i])
}

# Plotting histogram of prior crossings
hist(prior_crossings, breaks = 20, col = "skyblue", border = "white",
     xlab = "Number of Crossings", main = "Prior Predictions of Crossings")
```

Prior Predictions of Crossings



Exercise 2.3 -

Consider a dataset of crossing dependencies from English and German corpora, “crossing.csv”. This dataset contains number of crossings for each sentence from each language. Fit the following two models, M_1 and M_2 , to the given data.

Model M_1 :

- Assumption -

The rate of crossings is only a function of sentence length and it remains exactly the same in English and German.

- Regression - $N_{i,j} \sim \text{Poisson}(\lambda_{i,j})$ where $N_{i,j}$ is the number of crossing dependencies in sentence i in language j ; $\lambda_{i,j}$ is rate parameter indicating the expected rate of crossing dependencies in sentence i in language j , such that $\log(\lambda_{i,j}) = \alpha + \beta L_{i,j}$ where $L_{i,j}$ is the length of sentence i of language j .

Model M_2 :

- Assumption -

As sentence length increases, the number crossings grows at a different rate in English vs. German.

- Regression - $N_{i,j} \sim \text{Poisson}(\lambda_{i,j})$ where $N_{i,j}$ is the number of crossing dependencies in sentence i in language j ; $\lambda_{i,j}$ is rate parameter indicating the expected rate of crossing dependencies in sentence i in language j , such that $\log(\lambda_{i,j}) = \alpha + \beta L_{i,j} + \beta_{\text{language}} R_j + \beta_{\text{interact}} L_{i,j} * R_j$ where $L_{i,j}$ is the length of sentence i of language j , R_j is the indicator variable such that $R_j = 0$ if the language is English and $R_j = 1$ if the language is German.

Priors :

- $\alpha \sim \text{Normal}(0.15, 0.1)$
- $\beta \sim \text{Normal}(0, 0.15)$
- $\beta_{\text{language}} \sim \text{Normal}(0, 0.15)$
- $\beta_{\text{interact}} \sim \text{Normal}(0, 0.15)$

```
observed <- read.table("crossings.csv", header = T, sep = ",")
flextable(head(observed))
```

Language	s.id	s.length	nCross
German	1	2	0
German	2	2	1
German	3	2	0
German	4	2	0

Language	s.id	s.length	nCross
German	5	2	2
German	6	2	1

```
nrow(observed)
```

```
## [1] 1900
```

```
# Code/center the predictors
observed$s.length <- observed$s.length - mean(observed$s.length)
observed$lang <- ifelse(observed$Language=="German",1,0)
```

```
m1.fit <- brm(
  nCross ~ 1 + s.length,
  family = poisson(link = "log"),
  data = observed,
  prior = c(prior(normal(0.15, 0.1), class = Intercept),
            prior(normal(0, 0.15), class = b)),
  cores = 4,
  verbose = FALSE
)
```

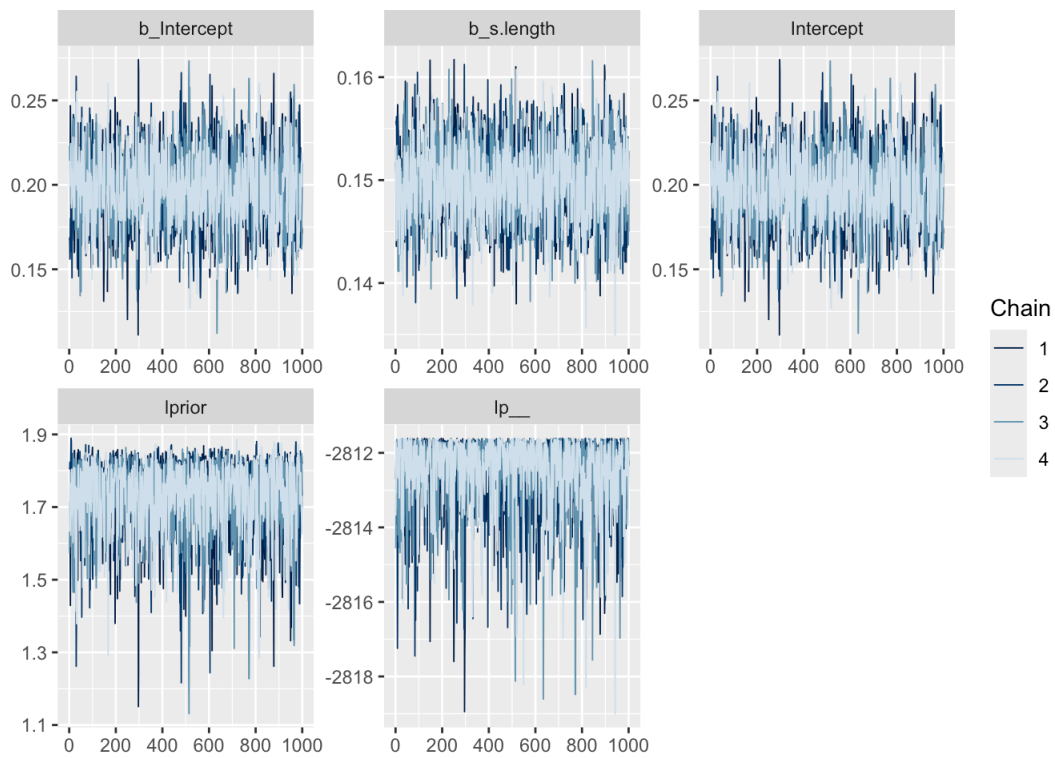
```
summary(m1.fit)
```

```
## Family: poisson
## Links: mu = log
## Formula: nCross ~ 1 + s.length
## Data: observed (Number of observations: 1900)
## Draws: 4 chains, each with iter = 2000; warmup = 1000; thin = 1;
## total post-warmup draws = 4000
##
## Regression Coefficients:
##           Estimate Est.Error l-95% CI u-95% CI Rhat Bulk_ESS Tail_ESS
## Intercept    0.20     0.02    0.15    0.24 1.00    1394    2009
## s.length     0.15     0.00    0.14    0.16 1.00    1627    2224
##
## Draws were sampled using sampling(NUTS). For each parameter, Bulk_ESS
## and Tail_ESS are effective sample size measures, and Rhat is the potential
## scale reduction factor on split chains (at convergence, Rhat = 1).
```

After fitting model M_1 using the data we get α - 0.1968321, 0.0218648, 0.1549425, 0.2393788


After fitting model M_1 using the data we get β - 0.1497753, 0.0036893, 0.1425566, 0.1568218

```
mcmc_trace(m1.fit)
```



```
mcmc_hist(m1.fit, pars = c("b_Intercept", "b_s.length"))
```

```
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```



```
m2.fit <- brm(
  nCross ~ 1 + s.length + lang + s.length*lang,
  family = poisson(link = "log"),
  data = observed,
  prior = c(prior(normal(0.15, 0.1), class = Intercept),
            prior(normal(0, 0.15), class = b)),
  cores = 4,
  verbose = FALSE
)
```

```
summary(m2.fit)
```

```
## Family: poisson
## Links: mu = log
## Formula: nCross ~ 1 + s.length + lang + s.length * lang
## Data: observed (Number of observations: 1900)
## Draws: 4 chains, each with iter = 2000; warmup = 1000; thin = 1;
##         total post-warmup draws = 4000
##
## Regression Coefficients:
##           Estimate Est.Error l-95% CI u-95% CI Rhat Bulk_ESS Tail_ESS
## Intercept      0.16      0.03    0.10    0.22 1.00    2058    2155
## s.length       0.10      0.01    0.09    0.11 1.00    2053    2048
## lang           0.02      0.05   -0.07    0.12 1.00    2033    2105
## s.length:lang   0.10      0.01    0.08    0.11 1.00    1976    2063
##
## Draws were sampled using sampling(NUTS). For each parameter, Bulk_ESS
## and Tail_ESS are effective sample size measures, and Rhat is the potential
## scale reduction factor on split chains (at convergence, Rhat = 1).
```

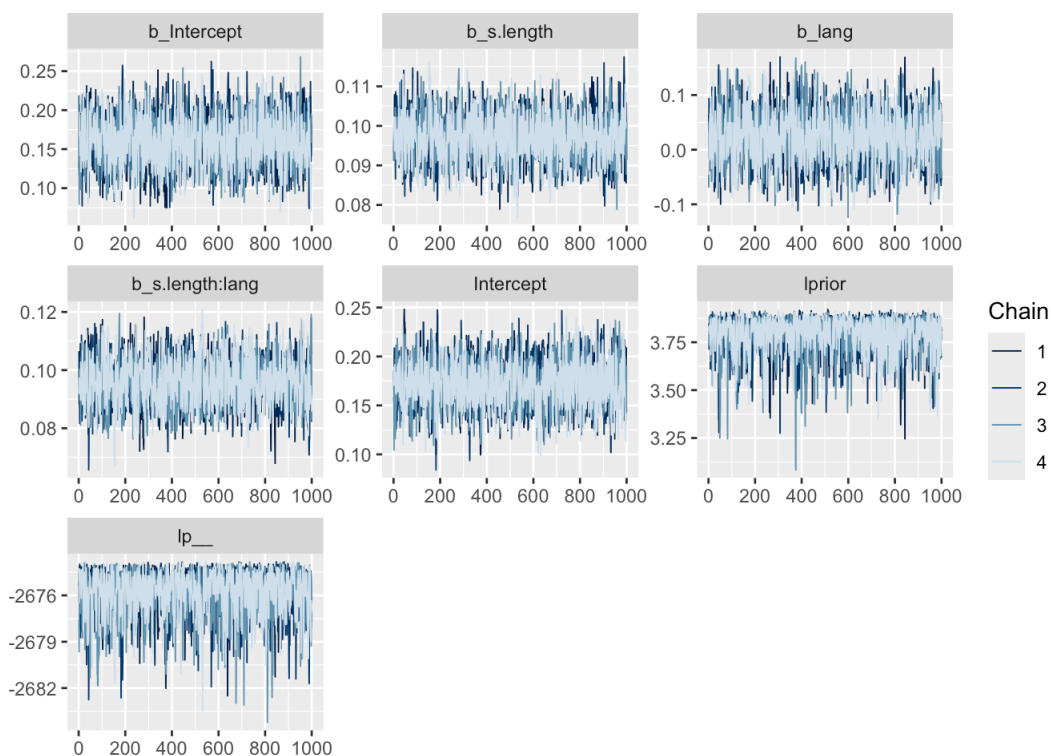
After fitting model M_2 using the data we get α - NA

After fitting model M_2 using the data we get β - NA

After fitting model M_2 using the data we get $\beta_{language}$ - NA


After fitting model M_2 using the data we get $\beta_{interaction}$ - NA

```
mcmc_trace(m2.fit)
```



```
mcmc_hist(m2.fit, pars = c("b_Intercept", "b_s.length", "b_lang", "b_s.length:lang"))
```

```
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```





Exercise 2.4 -

Quantify evidence for the models M1 and M2 using k-fold cross-validation.

```
observed %>% group_by(Language,s.length) %>%
  summarise(mean.crossings=mean(nCross)) %>%
  ggplot(aes(x=s.length,y=mean.crossings,
             group=Language,color=Language))+
  geom_point()+geom_line()
```

```
## `summarise()` has grouped output by 'Language'. You can override using the
## `.groups` argument.
```



```

# These two vectors will store log predictive densities
# in each fold
lpds.m1 <- c()
lpds.m2 <- c()

untested <- observed

for(k in 1:5){
  # Prepare test data and training data
  ytest <- sample_n(untested,size=nrow(observed)/5)
  ytrain <- setdiff(observed,ytest)

  untested <- setdiff(untested,ytest)
  # Fit the models M1 and M2 on training data
  fit.m1 <-
    brm(nCross ~ 1 + s.length,data=ytrain,
        family = poisson(link = "log"),
        prior = c(prior(normal(0.15, 0.1), class = Intercept),
                  prior(normal(0, 0.15), class = b)),
        cores=4,
        verbose = FALSE)
  fit.m2 <-
    brm(nCross ~ 1 + s.length + lang + s.length*lang,
        data=ytrain,
        family = poisson(link = "log"),
        prior = c(prior(normal(0.15, 0.1), class = Intercept),
                  prior(normal(0, 0.15), class = b)),
        cores=4,
        verbose = FALSE)
  # retrieve posterior samples
  post.m1 <- posterior_samples(fit.m1)
  post.m2 <- posterior_samples(fit.m2)

  # Calculate log pointwise predictive density using test data
  lppd.m1 <- 0
  lppd.m2 <- 0
  for(i in 1:nrow(ytest)){
    lpd_im1 <- log(mean(dpois(ytest[i,]$nCross,
                             lambda=exp(post.m1[,1]+
                             post.m1[,2]*ytest[i,]$s.length))))
    lppd.m1 <- lppd.m1 + lpd_im1
    lpd_im2 <- log(mean(dpois(ytest[i,]$nCross,
                             lambda=exp(post.m2[,1]+
                             post.m2[,2]*ytest[i,]$s.length+
                             post.m2[,3]*ytest[i,]$lang+
                             post.m2[,4]*ytest[i,]$s.length*ytest[i,]$lang)
                             )))
    lppd.m2 <- lppd.m2 + lpd_im2
  }
  lpds.m1 <- c(lpds.m1,lppd.m1)
  lpds.m2 <- c(lpds.m2,lppd.m2)
}
# Predictive accuracy of model M1
elpd.m1 <- sum(lpds.m1)
elpd.m1_SE <- sqrt(5*var(lpds.m1))

# Predictive accuracy of model M2
elpd.m2 <- sum(lpds.m2)
elpd.m2_SE <- sqrt(5*var(lpds.m2))

# Evidence in favor of M2 over M1
difference_elpd <- elpd.m2-elpd.m1

```

```
elpd.m1
```

```
## [1] -2814.18
```

```
elpd.m1_SE
```

```
## [1] 36.64415
```

```
elpd.m2
```

```
## [1] -2680.777
```

```
elpd.m2_SE
```

```
## [1] 13.41668
```

```
difference_elpd
```

```
## [1] 133.4034
```

Conclusions based on the above evidence :

- Model M_2 has a higher ELPD, indicating better predictive accuracy compared to Model M_1 .
- The positive difference (difference_elpd=134.1365) favors Model M_2 , suggesting that Model M_2 outperforms Model M_1 in terms of predictive accuracy.
- Both models have relatively low standard errors, indicating reasonably precise estimates of their respective ELPD values. However, Model M_2 has a slightly lower standard error, suggesting greater confidence in its ELPD estimate compared to Model M_1 .